

Text Preprocessing (II)

Text and Web Mining (H6751)

WKW School of Communication and Information


Word Sense Disambiguation

- English words are very often **ambiguous** as to their meaning or reference. [**Homonym, polysemy** – word with multiple meanings depending on context.
- For the example “**bore**,” one cannot tell without context even with POS - if the word is referring to a person—“**he is a bore**”—or a reference to a hole, as in “**the bore is not large enough.**”
 - “bore” is tagged as **NN** (Noun, singular or mass).
 - e.g., may have following features: bore:sense1(**person**), bore:sense2(**hole**)

Word Sense Disambiguation

- Word Sense Disambiguation – identifying the correct sense or semantics of a word based on its usage (**context**)

I killed the mouse
while working on my
report.



OR



- There are no algorithms that can completely disambiguate a text. Sometimes, humans can't too!
- How can we perform word sense disambiguation with limited context?

Word Sense Disambiguation

- Let's disambiguate “**bank**” in this sentence:

*“The **bank** can guarantee deposits will eventually cover future tuition costs because it invests in adjustable-rate mortgage securities.”*

- Given the following two WordNet senses:

bank ¹	Gloss:	a financial institution that accepts deposits and channels the money into lending activities
	Examples:	“he cashed a check at the bank”, “that bank holds the mortgage on my home”
bank ²	Gloss:	sloping land (especially the slope beside a body of water)
	Examples:	“they pulled the canoe up on the bank”, “he sat on the bank of the river and watched the currents”

The Simplified Lesk algorithm

Choose sense with **most word overlap** between gloss and context
(not counting function words – word that expresses grammatical or structural relationship with other words in a sentence)

“The **bank** can guarantee **deposits** will eventually cover future tuition costs because it invests in adjustable-rate **mortgage** securities.”

bank ¹	Gloss:	a financial institution that accepts deposits and channels the money into lending activities
	Examples:	“he cashed a check at the bank”, “that bank holds the mortgage on my home”
bank ²	Gloss:	sloping land (especially the slope beside a body of water)
	Examples:	“they pulled the canoe up on the bank”, “he sat on the bank of the river and watched the currents”

Word Sense Disambiguation

- Performs the classic Lesk algorithm for Word Sense Disambiguation (<http://www.nltk.org/howto/wsd.html>).
- Given an ambiguous word and the context in which the word occurs, Lesk returns a **Synset** with the highest number of overlapping words between the context sentence and different definitions from each Synset.

```
>>> from nltk.wsd import lesk
>>> sent = ['I', 'went', 'to', 'the', 'bank', 'to', 'deposit', 'money', '.']

>>> print(lesk(sent, 'bank', 'n'))
Synset('savings_bank.n.02')

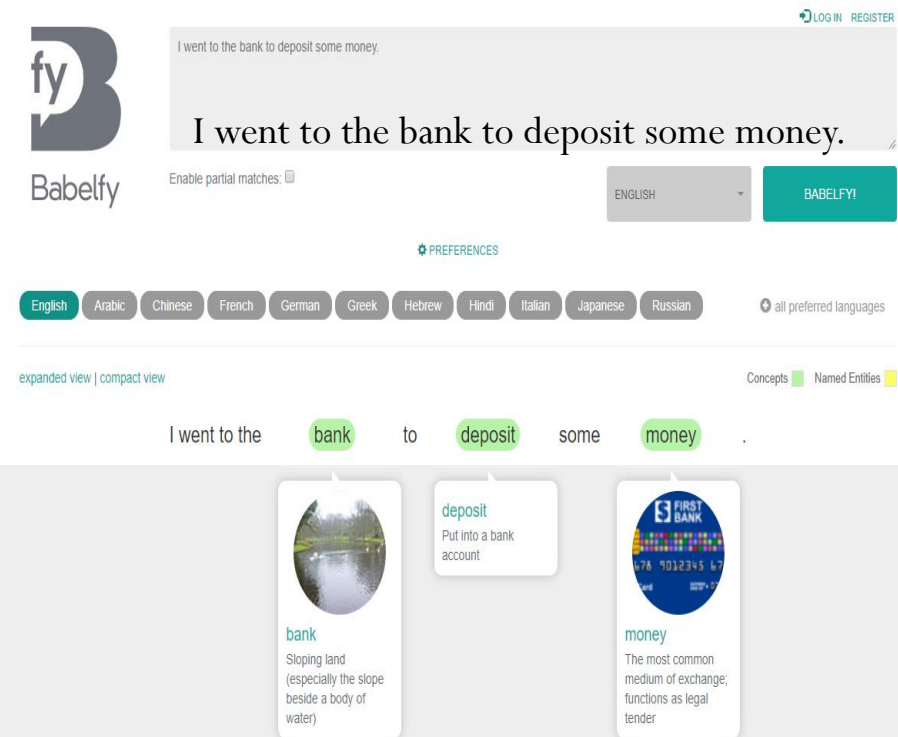
>>> print(lesk(sent, 'bank'))
Synset('savings_bank.n.02')
```

```
from nltk.corpus import wordnet as wn
for ss in wn.synsets('bank'):
    print(ss, ss.definition())
```

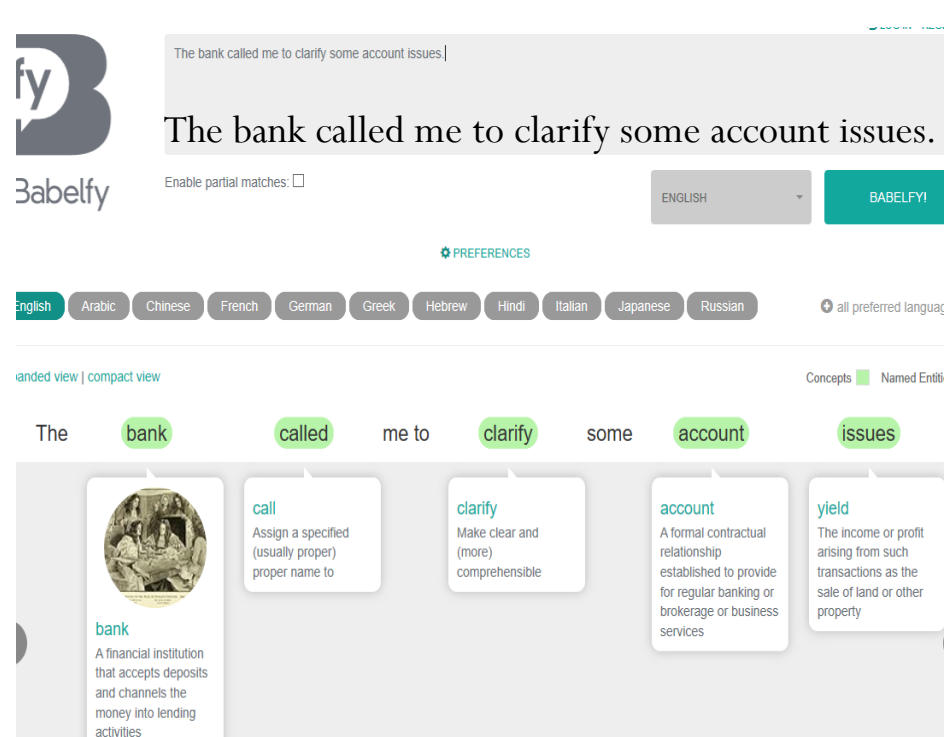
```
Synset('bank.n.01') sloping land (especially the slope beside a body of water)
Synset('depository_financial_institution.n.01') a financial institution that accepts deposits and channels the money into lending activities
Synset('bank.n.03') a long ridge or pile
Synset('bank.n.04') an arrangement of similar objects in a row or in tiers
Synset('bank.n.05') a supply or stock held in reserve for future use (especially in emergencies)
Synset('bank.n.06') the funds held by a gambling house or the dealer in some gambling games
Synset('bank.n.07') a slope in the turn of a road or track; the outside is higher than the inside in order to reduce the effects of centrifugal force
Synset('savings_bank.n.02') a container (usually with a slot in the top) for keeping money at home
Synset('bank.n.09') a building in which the business of banking transacted
Synset('bank.n.10') a flight maneuver; aircraft tips laterally about its longitudinal axis (especially in turning)
Synset('bank.v.01') tip laterally
Synset('bank.v.02') enclose with a bank
Synset('bank.v.03') do business with a bank or keep an account at a bank
Synset('bank.v.04') act as the banker in a game or in gambling
Synset('bank.v.05') be in the banking business
Synset('deposit.v.02') put into a bank account
Synset('bank.v.07') cover with ashes so to control the rate of burning
Synset('trust.v.01') have confidence or faith in
```

Word Sense Disambiguation

- Uses **Babelify** for Word Sense Disambiguation (WSD) (<http://babelify.org/>).
- Considered as a state-of-the-art system based on the BabelNet multilingual semantic network for multilingual Word Sense Disambiguation and Entity Linking.



The screenshot shows the Babelify web interface. At the top, the Babelify logo is on the left, and 'LOG IN' and 'REGISTER' links are on the right. Below the logo is the text 'Babelify'. The main input area contains the sentence 'I went to the bank to deposit some money.' Below this is a checkbox for 'Enable partial matches:' and a language dropdown set to 'ENGLISH'. A teal button labeled 'BABELFY!' is to the right. Below the input area is a 'PREFERENCES' link. A row of language buttons (English, Arabic, Chinese, French, German, Greek, Hebrew, Hindi, Italian, Japanese, Russian) and a link 'all preferred languages' are shown. Below this is a toggle for 'expanded view | compact view' and a legend for 'Concepts' (green square) and 'Named Entities' (yellow square). The sentence is displayed with words highlighted in green boxes: 'bank', 'deposit', and 'money'. Below each highlighted word is a card showing a representative image and a definition. For 'bank', the image shows a riverbank and the definition is 'Sloping land (especially the slope beside a body of water)'. For 'deposit', the image shows a bank account screen and the definition is 'Put into a bank account'. For 'money', the image shows a banknote and the definition is 'The most common medium of exchange; functions as legal tender'.



The screenshot shows the Babelify web interface for a different sentence. The main input area contains the sentence 'The bank called me to clarify some account issues.' Below this is a checkbox for 'Enable partial matches:' and a language dropdown set to 'ENGLISH'. A teal button labeled 'BABELFY!' is to the right. Below the input area is a 'PREFERENCES' link. A row of language buttons (English, Arabic, Chinese, French, German, Greek, Hebrew, Hindi, Italian, Japanese, Russian) and a link 'all preferred languages' are shown. Below this is a toggle for 'expanded view | compact view' and a legend for 'Concepts' (green square) and 'Named Entities' (yellow square). The sentence is displayed with words highlighted in green boxes: 'bank', 'called', 'clarify', 'account', and 'issues'. Below each highlighted word is a card showing a representative image and a definition. For 'bank', the image shows a bank building and the definition is 'A financial institution that accepts deposits and channels the money into lending activities'. For 'called', the image shows a telephone and the definition is 'Assign a specified (usually proper) proper name to'. For 'clarify', the image shows a person writing and the definition is 'Make clear and (more) comprehensible'. For 'account', the image shows a bank document and the definition is 'A formal contractual relationship established to provide for regular banking or brokerage or business services'. For 'issues', the image shows a document and the definition is 'The income or profit arising from such transactions as the sale of land or other property'.

- Unless a project (e.g. Q&A system, Information Retrieval) requires word sense disambiguation, it is best to proceed without such a step.

BeautifulSoup

- It's difficult to decipher textual content in web page due to unnecessary HTML tags.

```
sample = """<h1>Title Goes Here</h1>

<b>Bolded Text</b>
<i>Italicized Text</i>


<a href="this will be gone, too">But this will still be here!</a>

I run. He ran. She is running. Will they stop running?
I talked. She was talking. They talked to them about running. Who ran
```

- The BeautifulSoup library provides functions to remove tags with ease.

```
def strip_html(text):
    soup = BeautifulSoup(text, "html.parser")
    return soup.get_text()

sample = strip_html(sample)
print(sample)
```

Title Goes Here
Bolded Text
Italicized Text

But this will still be here!

I run. He ran. She is running. Will they stop running?
I talked. She was talking. They talked to them about running. Who ran to the talking runner?

Expanding Contractions

- Contractions are shortened versions of words. Avoided in formal writing but used quite extensively in informal communications.
- Problem for text mining – contractions represent same meaning of contracted words. Need for normalization.
- **Contraction Maps** to match corresponding versions of words. “**Contractions**” library available in Python.

```
CONTRACTION_MAP = {  
    "ain't": "is not",  
    "aren't": "are not",  
    "can't": "cannot",  
    "can't've": "cannot have",  
    .  
    .  
    .  
    "you'll've": "you will have",  
    "you're": "you are",  
    "you've": "you have"
```

```
!pip install contractions
```

```
import contractions
```

```
text = "I can't go to the movies. We don't want to buy the books."
```

```
sample = contractions.fix(text) # e.g., can't -> cannot; don't -> do not  
print(sample)
```

I can not go to the movies. We do not want to buy the books.

Regular expressions



- Regular expressions (aka regexes) create string patterns and use them for searching and substituting specific pattern matches in text.
 - How to search for Woodchuck, woodchuck, Woodchucks, woodchucks

Operator	Function
^	matches the start of the string. E.g., ^The matches string that start with The
\$	matches the end of the string. E.g., end\$ matches string that ends with end
Braces { }	Indicate range of preceding occurrences. E.g., ab{2} => abb , ab{2,3} => abb , abbb
Square Bracket [...]	matches any one of the set of characters in []. E.g., [ab] matches a or b. Caret ^ appears first in [] negates pattern. [^ab] => matches anything except a or b
Parentheses ()	captures all the characters within () as a group. E.g., (the) matches “ the ” in “ then ”

Pattern	Matches
[wW]oodchuck	Woodchuck, woodchuck
(456)	456 in 1234 56 7890

Regular expressions

Representation Operators

Operator	Function
.	matches a single character. E.g., a.e matches ate , a1e , a@e
\w	matches a letter or digit or underbar => [a-zA-Z0-9_]. \W matches any non \w
\s	matches a single whitespace character. \S matches non-whitespace character.
\d	matches a single digit => [0-9]

Repetition Operators

Operator	Function
*	matches zero or more cases of the previous mentioned regex before the * symbol in the pattern. E.g., ab* matches “ac”, “abc”, “abbc”
+	matches one or more cases of the previous mentioned regex before the + symbol in the pattern. E.g., ab+ matches “ab”, “abbc” but not “ac”
?	matches zero or one case of the previous mentioned regex before the ? symbol in the pattern. E.g., ab? matches “ac”, “abc” but not “abbc”
 	OR operator. E.g., a b matches a or b

Regular Expressions

- The “**re**” Python module is used for string searching and manipulation
 - **re.match(pattern, string)**: This method is used to match pattern at the **beginning** of string.

```
>> re.match('C', 'IceCream').group()
>> None
```

```
>> re.match('C', 'Cake').group()
>> C
```

- **re.search()**: This method is used to match patterns occurring at **any position** in the string.

```
>> re.search('cookie', 'Cake and cookie').group()
>> 'cookie'
```

- **re.findall()**: This method returns all non-overlapping matches of the specified regex pattern in the string.

```
email_address = "Please contact us at: support@abc.com, sales@abc.com"
>> results = re.findall('[\w]+@[ \w.] +', email_address)
```

matches letter

Returns ["support@abc.com", sales@abc.com].

Part-of-Speech Tagging

- If the text mining goal is specific, say recognizing names of **people**, **places**, and **organizations**, it is usually desirable to perform additional linguistic analyses of the text and extract more sophisticated features.
 - E.g., San Francisco: San/NNP Francisco/NNP (NNP: Proper noun, singular)
- In English, some analyses may use as few as six or seven categories and others nearly one hundred.
- Most English grammars would have as a minimum **noun**, **verb**, **adjective**, **adverb**, **preposition**, and **conjunction**.
- POS can be used for **feature reduction**, e.g., use only verb, adjective, and adverb for sentiment classification.
- Distribution of POS can be used for author, gender, and document genre (formal vs. informal) classification.

Part-of-Speech Tagging

- A set of 36 categories, constructed from the Wall Street Journal corpus is used in the **Penn Tree Bank**
- A **tree bank** is a parsed text corpus that annotates sentence structure, such as POS and phrases.
- Almost all POS taggers have been trained on the Wall Street Journal corpus available from LDC (Linguistic Data Consortium, www.ldc.upenn.edu)

Number	Tag	Description			
1.	CC	Coordinating conjunction	19.	PRP\$	Possessive pronoun
2.	CD	Cardinal number	20.	RB	Adverb
3.	DT	Determiner	21.	RBR	Adverb, comparative
4.	EX	Existential <i>there</i>	22.	RBS	Adverb, superlative
5.	FW	Foreign word	23.	RP	Particle
6.	IN	Preposition or subordinating conjunction	24.	SYM	Symbol
7.	JJ	Adjective	25.	TO	<i>to</i>
8.	JJR	Adjective, comparative	26.	UH	Interjection
9.	JJS	Adjective, superlative	27.	VB	Verb, base form
10.	LS	List item marker	28.	VBD	Verb, past tense
11.	MD	Modal	29.	VBG	Verb, gerund or present participle
12.	NN	Noun, singular or mass	30.	VCN	Verb, past participle
13.	NNS	Noun, plural	31.	VBP	Verb, non-3rd person singular present
14.	NNP	Proper noun, singular	32.	VBZ	Verb, 3rd person singular present
15.	NNPS	Proper noun, plural	33.	WDT	Wh-determiner
16.	PDT	Predeterminer	34.	WP	Wh-pronoun
17.	POS	Possessive ending	35.	WPS	Possessive wh-pronoun
18.	PRP	Personal pronoun	36.	WRB	Wh-adverb

Part-of-Speech Tagging

- **The Stanford Parser:** a statistical parser

An implementation in Java: <http://nlp.stanford.edu/software/lex-parser.shtml>

Stanford Parser

Please enter a sentence to be parsed:

My dog also likes eating sausage.

Language: English ▾

Sample Sentence

Parse

Your query

My dog also likes eating sausage.

Tagging

My/PRP\$ dog/NN also/RB likes/VBZ eating/VBG sausage/NN ./.

- **Online Stanford Parser**

<http://nlp.stanford.edu:8080/parser/>

Note:

PRP\$: Possessive pronoun

RB: Adverb

VBZ: Verb, 3rd person singular present

Phrase Recognition

- Once the tokens of a sentence have been assigned **POS tags**, the next step is to group individual tokens (one being the main or head word) into units, generally called **phrases**.
- **Noun phrases** – act as a subject or object to a verb. Consist of a **noun** or pronoun [**head word**], and **dependent** words before or after the head. [e.g., “the moon”]
- **Verb phrases** – consist of a main verb [MV] alone, or a main verb plus any modal [MO] and/or auxiliary [AUX] verbs.

[MV]

We all laughed.

[MO] [AUX][AUX] [MV]

Tony might have been waiting outside for you

- **Prepositional phrases** – consists of a preposition [head] (e.g., to, with), its object (noun/pronoun), and object modifiers (article/adjective). Article – determiner that precedes a noun.

e.g., The boy *with her* is her son

Phrase Recognition

- **The Stanford Parser: online parser**

<http://nlp.stanford.edu:8080/parser/>

Please enter a sentence to be parsed:

My dog also likes eating sausage.

Language: English ▼

Sample Sentence

Parse

Your query

My dog also likes eating sausage.

Tagging





My/PRP\$ dog/NN also/RB likes/VBZ eating/VBG sausage/NN ./.

Parse

```
(ROOT
  (S
    (NP (PRP$ My) (NN dog))
    (ADVP (RB also))
    (VP (VBZ likes)
      (S
        (VP (VBG eating)
          (NP (NN sausage)))))
    (. .)))
```

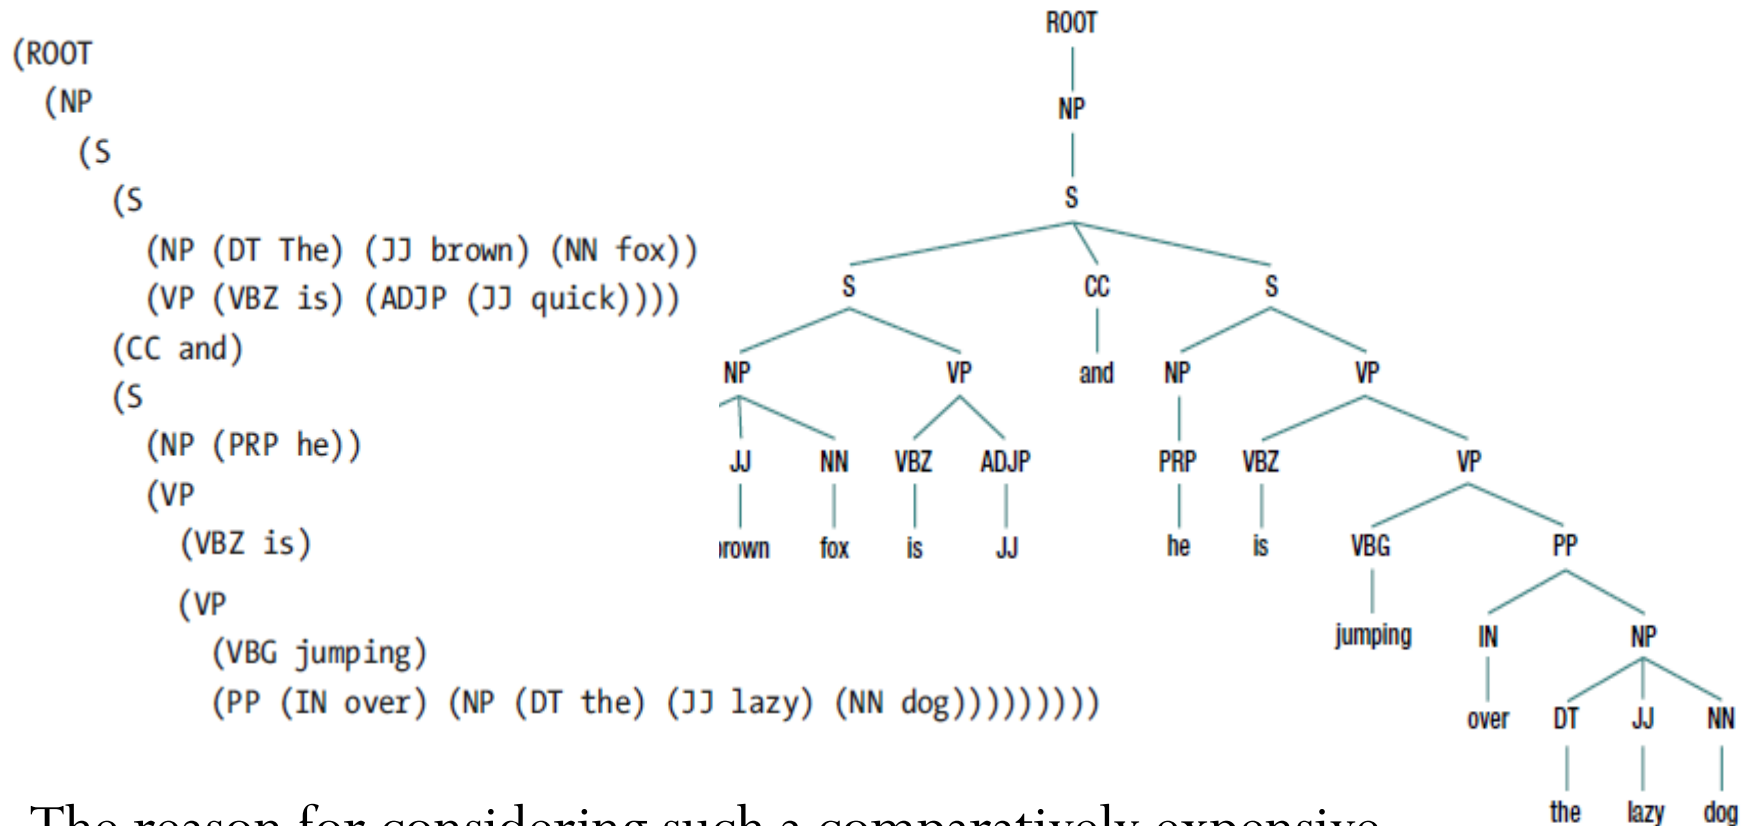
Parsing (Phrase Structure Rules)

- Parsing - step of producing a **full parse of a sentence** based on Phrase Structure Rules
 - Generic rule denotes binary division for a sentence or a clause [$S \rightarrow AB$], structure S consists of constituents A and B with the order A followed by B.
 - $S \rightarrow NP VP \Rightarrow$ sentence or clause divided into the subject (NP) and predicate (VP)
 - $NP \rightarrow [DET][ADJ]N[PP] \Rightarrow$ Noun as the head word optional Determinants, Adjectives, and Prepositional
 - $VP \rightarrow V [VP][NP][PP][ADJP][ADVP] \Rightarrow$ Verb head word followed by optional VP, NP, PP, Adjective Phrase or Adverbial Phrase.
 - $PP \rightarrow PREP [NP] \Rightarrow$ preposition as head word followed by optional NP



Parsing (Phrase Structure Rules)

Sentence – The brown fox is quick and he is jumping over the lazy dog.



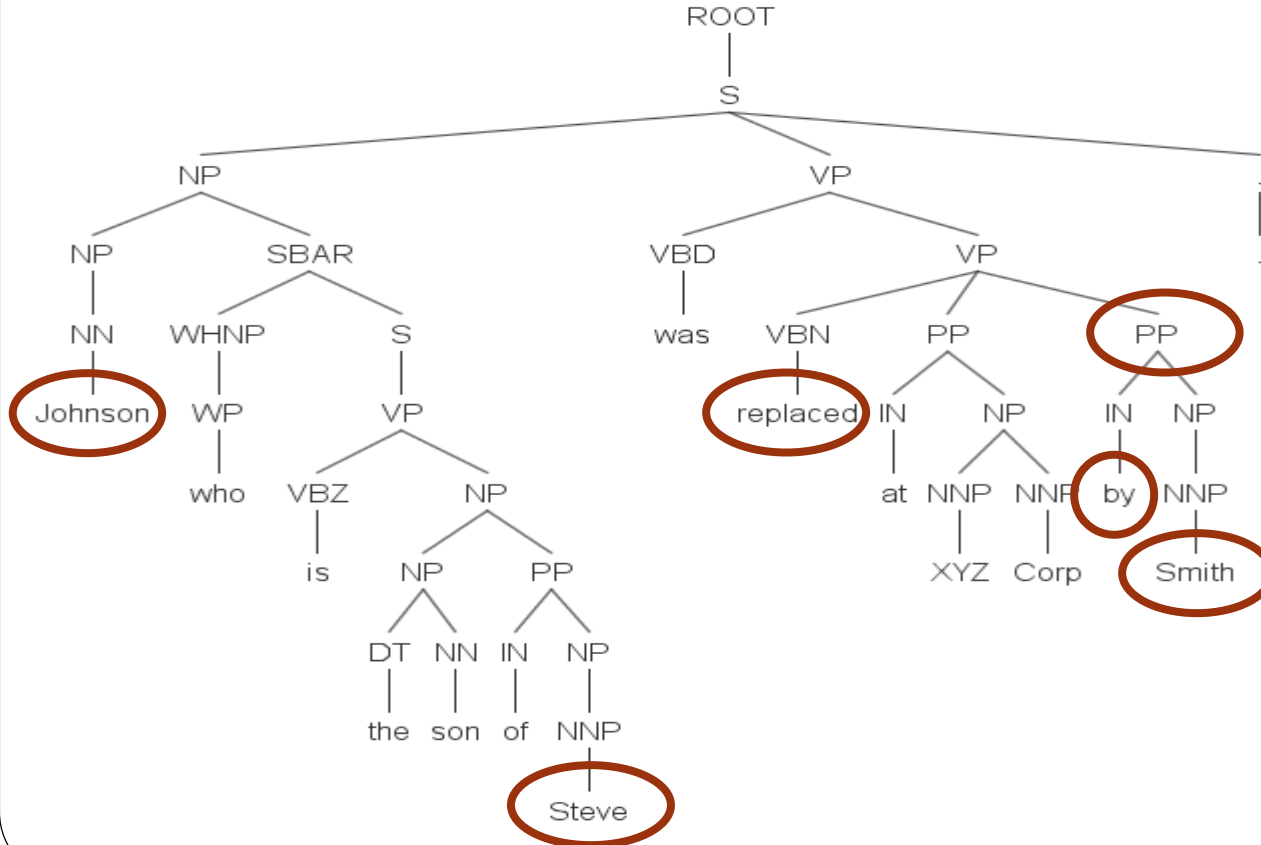
- The reason for considering such a comparatively expensive process is that it provides **detailed syntactic relationships information** that phrase identification cannot provide.

Parsing

- Consider the sentence “Johnson who is the son of Steve was replaced at XYZ Corp by Smith.”

Using the Parse Tree, machine can infer that Johnson was replaced by Smith (correct) and not Steve was replaced by Smith (wrong).


Phrase structure tree



```
(ROOT
  (S
    (NP
      (NP (NNP Johnson))
      (SBAR
        (WHNP (WP who))
        (S
          (VP (VBZ is)
            (NP
              (NP (DT the) (NN son))
              (PP (IN of)
                (NP (NNP Steve)))))))
          (VP (VBD was)
            (VP (VBN replaced)
              (PP (IN at)
                (NP (NNP XYZ) (NNP Corp)))
              (PP (IN by)
                (NP (NNP Smith))))))
          (. .)))
```

Parsing - how to process a Parse Tree?

- **Universal dependencies** (i.e. grammatical relations; evolved out of Stanford Dependencies) from [Stanford Parser](#): “I like him.”




Universal dependencies

```
nsubj(like-2, I-1)
root(ROOT-0, like-2)
dobj(like-2, him-3)
```

Output from [Stanford Parser](#)


Basic Dependencies:



Output from [Stanford CoreNLP](#)

Parsing

- Universal dependencies (i.e. grammatical relations) from [Stanford CoreNLP](#): “Johnson who is the son of Steve was replaced at XYZ Corp by Smith.”



```
(ROOT
  (S
    (NP
      (NP (NNP Johnson))
      (SBAR
        (WHNP (WP who))
        (S
          (VP (VBZ is)
            (NP
              (NP (DT the) (NN son))
              (PP (IN of)
                (NP (NNP Steve)))))))
          (VP (VBD was)
            (VP (VBN replaced)
              (PP (IN at)
                (NP (NNP XYZ) (NNP Corp)))
              (PP (IN by)
                (NP (NNP Smith))))))
          (. .)))
```

```
nsubjpass(replaced-9, Johnson-1)
nsubj(son-5, who-2)
cop(son-5, is-3)
det(son-5, the-4)
act:relcl(Johnson-1, son-5)
case(Steve-7, of-6)
nmod(son-5, Steve-7)
auxpass(replaced-9, was-8)
root(ROOT-0, replaced-9)
case(Corp-12, at-10)
compound(Corp-12, XYZ-11)
nmod(replaced-9, Corp-12)
case(Smith-14, by-13)
nmod(replaced-9, Smith-14)
```

Vector Generation for Prediction

- Without any deep analysis of the linguistic content of the documents, we can describe **each document by features that represent the most frequent tokens**.
- Each row is a document, and each column represents a feature.
- Thus, a cell in the spreadsheet is a measurement of a feature (corresponding to the column) for a document (corresponding to a row).

DocID	Apple	Bear	Durian	...	Zoo	Animal?
1	0	3	0	0	2	1
2	1	0	2	0	0	0
...						

- Dictionary (or feature) reduction techniques
 - Local dictionary, removing Stopwords, Frequent words, Feature selection, and Token reduction (stemming and synonyms)

Bag of Words Model

- A vector space model represents unstructured text (or any other data) as numeric vectors, such that each dimension of the vector is a specific feature/attribute.
- BOW represents each document as a numeric vector where each dimension is a specific word from the corpus and the value could be its frequency in the document, occurrence (denoted by 1 or 0), or even weighted values.
- Represented literally as a bag of its own words, disregarding word order, sequences, and grammar.

Bag of Words Example

Document 1

The quick brown fox jumped over the lazy dog's back.

Document 2

Now is the time for all good men to come to the aid of their party.

Term	Document 1	Document 2
	Document 1	Document 2
aid	0	1
all	0	1
back	1	0
brown	1	0
come	0	1
dog	1	0
fox	1	0
good	0	1
jump	1	0
lazy	1	0
men	0	1
now	0	1
over	1	0
party	0	1
quick	1	0
their	0	1
time	0	1

Stopword
List

for
is
of
the
to

CountVectorizer

- The CountVectorizer class can **produce a bag-of-words representation** from a string or file. (https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html)

```
class sklearn.feature_extraction.text.CountVectorizer(input='content', encoding='utf8', decode_error='strict', strip_accents=None, lowercase=True, preprocessor=None, tokenizer=None, stop_words=None, token_pattern='(?u)\b\w\w+\b', ngram_range=(1, 1), analyzer='word', max_df=1.0, min_df=1, max_features=None, vocabulary=None, binary=False, dtype=<class 'numpy.int64'>)
```

- By default, CountVectorizer converts the characters in the documents to **lowercase**, and **tokenizes** the documents using a regular expression that splits strings on whitespace and extracts sequences of characters that are two or more characters in length.

CountVectorizer

- The documents in the corpus can be represented by **feature vectors**:

```
>>> from sklearn.feature_extraction.text import CountVectorizer
```

```
>>> corpus = [
```

```
>>> 'UNC played Duke in basketball',
```

```
>>> 'Duke lost the basketball game'
```

```
>>> ]
```

```
>>> vectorizer = CountVectorizer()
```

```
>>> print vectorizer.fit_transform(corpus).todense()
```

```
>>> print vectorizer.vocabulary_
```

```
[[1 1 0 1 0 1 0 1]
```

```
[1 1 1 0 1 0 1 0]]
```

```
{u'duke': 1, u'basketball': 0, u'lost': 4, u'played': 5, u'game': 2, u'unc': 7,  
u'in': 3, u'the': 6}
```

	basketball	duke	game	in	lost	played	the	unc
D1	1	1	0	1	0	1	0	1
D2	1	1	1	0	1	0	1	0

Bag of N-Grams Model

- A word is just a single token, often known as a *unigram* or *1-gram*. Bag of Words model doesn't consider the order of words. But what if we also wanted to take into account phrases or collection of words that occur in a sequence?
- An N-gram is basically a collection of word tokens from a text document such that these tokens are contiguous and occur in a sequence.
- Bi-gram indicates n-grams of order 2 (two words e.g., [beautiful sky], [sky today]), tri-grams order 3 (three words e.g., [beautiful sky today]), and so on.

TF-IDF Model

- Potential problems with BOW model – words (terms) that occur frequently across all documents tend to overshadow other terms.
- **Bag-of-words** feature vectors do not encode **grammar**, **word order**, or **frequencies** of words. The frequency with which a word appears in a document could indicate the extent to which the document relates to the word (**term weights**).
- Especially words that don't occur as frequently, but might be interesting and effective as features.

Term Frequency

- **Term frequency (TF)** – raw frequency of a term in document $tf(t, d)$
- Words might appear with the same frequency in two documents, but the documents could be dissimilar if one is **many times larger** than the other.
 - “I love **apple**” vs “I love fruits especially **apples, oranges, pears, ...**”
- Mitigate problem by using **normalized** term frequency weights.

$$tf(t, d) = \frac{f(t, d)}{\|x\|} \quad \text{Term vector for } d = tf(\text{“I”}) + tf(\text{“love”}) + tf(\text{“apple”})$$

$f(t, d)$ is the frequency of term in document d and x is the L2 norm of the count vector of terms in the document $\Rightarrow \sqrt{x_1^2 + x_2^2 + x_3^2 + \dots}$

Term Frequency-Inverse Document Frequency

- **TF-IDF** is used instead of the raw frequencies of a term to **scale down impact of frequently occurring terms** in a given corpus, which are **less informative** than features that occur in a small fraction of the corpus.
- **Inverse Document Frequency (IDF)** measures how rare or common a word is in a corpus => more rare more significant.

Standard idf definition

- (i) $idf(d, t) = \log \left[\frac{(n)}{(df(d, t))} \right]$

Effect of adding “1” to idf is that terms with zero idf, i.e., terms that occur in all documents ($df=n$) leading to $\log(1)=0$, will not be entirely ignored.

- (ii) $idf(d, t) = \log \left[\frac{(1+n)}{(1+df(d, t))} \right] + 1$

Modified idf definition

Effect of adding “1” to df is that terms with zero df, i.e., terms that occur in no documents ($df=0$) will not lead to division by zero.

- **N** is the total number of documents in the corpus and $df(d, t)$ is the number of documents in the corpus that contain the term **t**. Log is used when `sublinear_tf=True` (**Default is LN**)

Term Frequency-Inverse Document Frequency

- A term's **TF-IDF** value is the product of its term frequency and inverse document frequency: $TF\text{-}IDF = TF \times IDF$.

$$\begin{aligned} TF\text{-}IDF &= tf(t, d) \cdot idf(d, t) \\ &= \frac{f(t, d)}{||x||} \cdot \log \left[\frac{(1+n)}{(1+df(d, t))} \right] + 1 \end{aligned}$$

- The resulting tf-idf vectors are then normalized by the Euclidean norm based on the tf-idf score of each term:

$$v_{norm} = \frac{v}{||v||_2} = \frac{v}{\sqrt{v_1^2 + v_2^2 + \dots + v_n^2}}$$

CountVectorizer and TfidfTransformer

- CountVectorizer - creation of feature vectors that encode the frequencies of words (term frequencies)

```
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
corpus = [
    'The dog ate a sandwich and I ate a sandwich',
    'The wizard transfigured a sandwich'
]
vectorizer = CountVectorizer(stop_words='english')
transformer = TfidfTransformer(use_idf=False)
transformerIDF = TfidfTransformer(use_idf=True)
X = vectorizer.fit_transform(corpus)
print('Count vectors:\n', X.todense())
print('Vocabulary:\n', vectorizer.vocabulary_)
print('TF vectors:\n', transformer.fit_transform(X).todense())
print('TF-IDF vectors:\n', transformerIDF.fit_transform(X).todense())
```

```
Count vectors:
[[2 1 2 0 0]
 [0 0 1 1 1]]
Vocabulary:
{'dog': 1, 'ate': 0, 'sandwich': 2, 'wizard': 4, 'transfigured': 3}
TF vectors:
[[ 0.66666667  0.33333333  0.66666667  0.          0.          ]
 [ 0.          0.          0.57735027  0.57735027  0.57735027]]
TF-IDF vectors:
[[ 0.75458397  0.37729199  0.53689271  0.          0.          ]
 [ 0.          0.          0.44943642  0.6316672  0.6316672 ]]
```

- TfidfTransformer** returns **TF-IDF weight** when the **use_idf** keyword argument is True (default value) (returns **TF weight** if use_idf set to False)

TF, IDF, and TF-IDF calculations

$$[\text{ate}, \text{doc1}] \text{ tf} = 2 / \sqrt{2^2 + 1^2 + 2^2} = 2 / \sqrt{9} = \mathbf{0.667}$$

Count vectors:

```
[[2 1 2 0 0]
 [0 0 1 1 1]]
```

$$[\text{wizard}, \text{doc2}] \text{ tf} = 1 / \sqrt{1^2 + 1^2 + 1^2} = 1 / \sqrt{3} = \mathbf{0.55735}$$

Vocabulary:

```
{'dog': 1, 'ate': 0, 'sandwich': 2, 'wizard': 4, 'transfigured': 3}
```

TF vectors:

```
[[0.66666667 0.33333333 0.66666667 0.          0.          ]
 [0.          0.          0.57735027 0.57735027 0.57735027]]
```

TF-IDF vectors:

```
[[0.75458397 0.37729199 0.53689271 0.          0.          ]
 [0.          0.          0.44943642 0.6316672  0.6316672  ]]
```

$$[\text{ate}, \text{doc1}] \text{ tfidf} = \text{tf} * \textcolor{red}{idf} = 0.667 * 1.405 = \mathbf{0.937}$$

$$\text{normalized tfidf} = \frac{0.937}{\sqrt{0.937^2 + 0.468^2 + 0.667^2}} = \mathbf{0.754}$$

$$\begin{aligned} \textcolor{red}{idf}(d, t) &= \ln \left[\frac{(1 + n)}{(1 + df(d, t))} \right] + 1 \\ &= \ln \left[\frac{(1 + 2)}{(1 + 1)} \right] + 1 \\ &= \mathbf{1.405} \end{aligned}$$

TfidfVectorizer

- Scikit-learn provides a **TfidfVectorizer** class that wraps **CountVectorizer** and **TfidfTransformer**.

```
>>> from sklearn.feature_extraction.text import TfidfVectorizer
>>> corpus = [
>>>     'The dog ate a sandwich and I ate a sandwich',
>>>     'The wizard transfigured a sandwich'
>>> ]
>>> vectorizer = TfidfVectorizer(stop_words='english')
>>> print vectorizer.fit_transform(corpus).todense()
```

Computes TF-IDF directly

0.75458397	0.37729199	0.53689272	0.	0.	0.
0.	0.	0.44943642	0.6316672	0.6316672	0.6316672

ate sandwich

- By comparing the TF-IDF weights to the raw term frequencies, we can see that words that are common to many of the documents in the corpus, such as **sandwich**, have been penalized.

Referenced Materials

- Fundamentals of Predictive Text Mining, Sholom M. Weiss, Nitin Indurkha, and Tong Zhang, Springer.
 - Chapter 2
- Natural Language Processing, Dan Jurafsky and Christopher Manning,
<http://www.stanford.edu/~jurafsky/NLPCourseraSlides.html>
- Machine Learning, Tom M. Mitchell, McGraw-Hill